

COPENHAGEN BUSINESS ACADEMY



DevOps part 3/4: Scaling & deployment

Jens Egholm Pedersen
<jeep@cphbusiness.dk>

Learning how to learn

- Meta-cognition
 - Think about how you think
 - Dunning-Kruger effect
- Use what we have prepared for you
- Continuous feedback
 - Where would you like to be when this course ends?
 - Keep evaluating yourself
 - ... and be honest!

See also: [Dunning-Kruger effect](#), [Metacognition improves your grade!](#)

Dev + Ops

- We expect you to develop your system
- And we expect you to keep an eye on them
- Do this for your own sake
 - You would like to learn about operating systems now
 - **Not** when you crash your production system for the first time
 - Trust me on this

Recap

- Service-level agreement (SLA)
- Monitoring
- Logging
- Post-mortem analysis

A word on bottlenecks

- All design problems can be solved by a layer of indirection
- No performance problem can be solved by removing a layer of indirection
- Bottlenecks appear in any good design
- Three problems:
 - Single point of entry
 - Single point of failure
 - Congestion

See also: [Profiling in software](#)

The problem with bottlenecks

- Single point of entry
 - How can we solve this?
- Single point of failure
 - How can we solve this?
- Congestion
 - How can we solve this?

Scaling the bottlenecks

- Single point of entry
 - Content delivery network
- Single point of failure
 - Redundancy
- Congestion
 - Scaling!

Kubernetes

- “Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications” - Google
- Powerful
- Complex
- We won't be using this

See also: [Docker swarm and Kubernetes compared](#)

Docker swarm mode

- Based on swarmkit
 - “A toolkit for orchestrating distributed systems at any scale”
- Orchestration
 - Manages nodes and services
- Scheduling
 - Resource aware task scheduling
- Security
 - We’ll talk about this next week

See also: [swarmkit on GitHub](#)

Docker commands

- `docker-machine`
 - Manages Vms
- `docker swarm`
 - Manages our cluster (swarm)
- `docker service`
 - Manages the containers in the swarm

Docker swarm: Concepts

- Nodes
 - A VM instance participating in a swarm
 - Worker nodes
 - Manager nodes
- Services and tasks
 - A collection of tasks to be executed on a node (= container)
 - Replicated services (fixed instances)
 - Global services (one instance per node)
- Load balancing

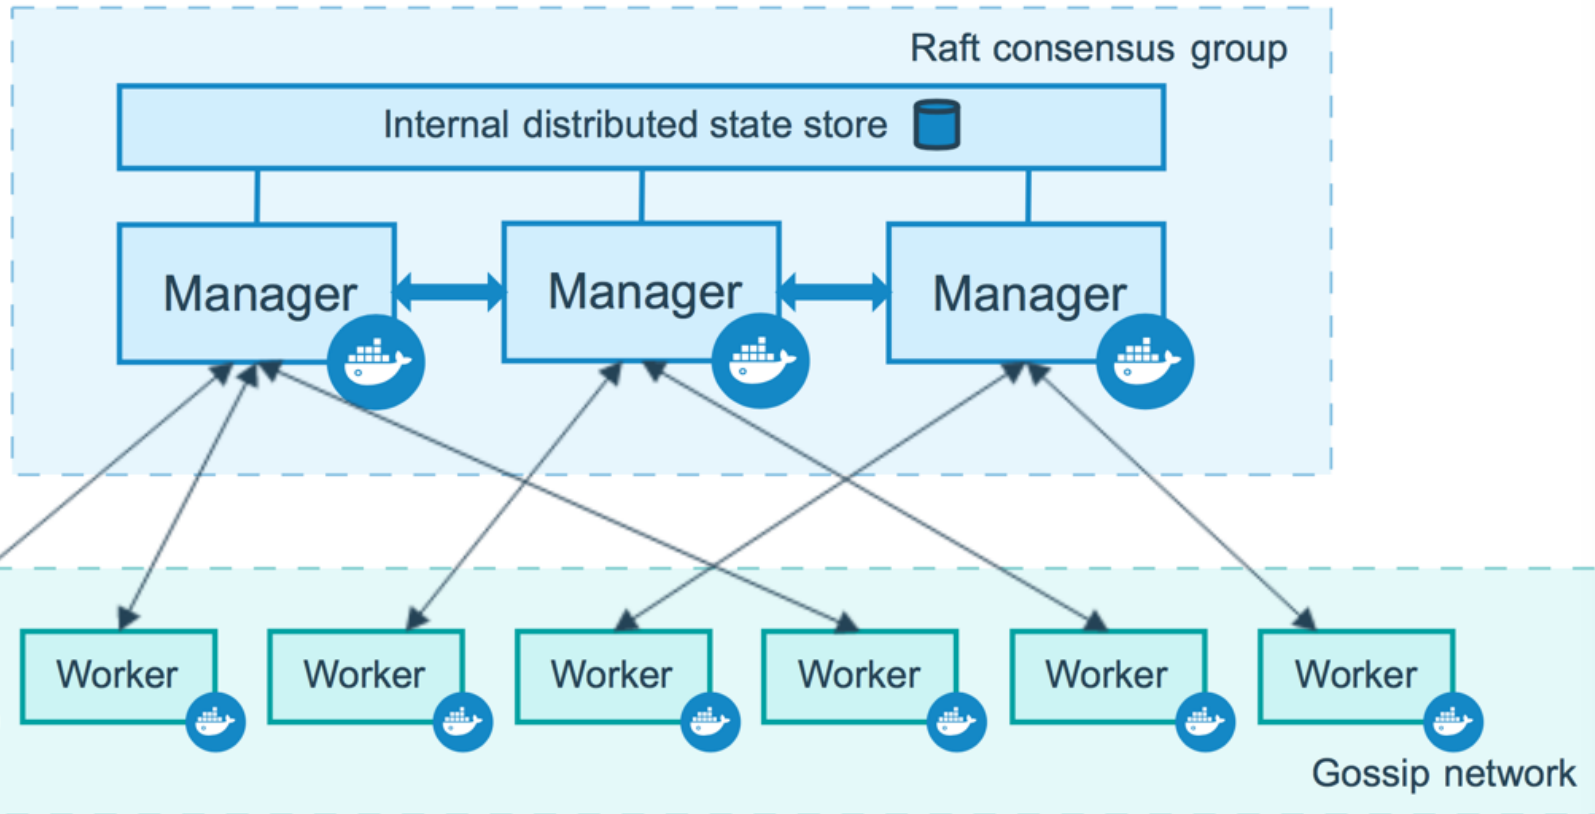
See also: [Swarm mode key concepts](#)

Docker swarm: Nodes

- Worker nodes
 - Contain a number of services
 - Need at least one manager node
- Manager node
 - Maintains cluster state and schedules services

See also: [Swarm nodes](#)

Docker swarm: Nodes



See also: [Swarm mode key concepts](#)

Introduction to Docker machine

- Can create and manage virtual machines for us
- Input parameters
 - Driver
 - Virtualbox, digitalocean, aws, etc.
 - Container name
 - Optional: Driver parameters
 - DO access token, DO region, etc.

See also: [Docker machine documentation](#)

Introduction to Docker machine

- Can create and manage virtual machines for us

- Locally with virtualbox

```
docker-machine create --driver virtualbox mybox
```

- Remotely with Digital Ocean

```
docker-machine create --driver digitalocean  
--digitalocean-access-token=xxx  
--digitalocean-region ams3 mybox
```

See also: [Docker machine documentation](#)

Introduction to Docker machine

- Can create and manage virtual machines for us

- Listing the available machines

```
docker-machine ls
```

- Getting the ip of a machine

```
docker-machine ip mybox
```

- Removing a machine

```
docker-machine rm mybox
```

- Ssh'ing into the available machines

```
docker-machine ssh mybox
```

See also: [Docker machine documentation](#)

Docker swarm: Managers

- Manager node
 - Maintains cluster state and schedules services
- Creating a swarm
 - On your swarm manager!!
`docker swarm init --advertise-addr x.x.x.x`
- Now we have a cluster!
 - With one node...

See also: [Creating a swarm](#)

Docker swarm: Managers

- Fault tolerance

Swarm Size	Majority	Fault Tolerance
1	1	0
2	2	0
3	2	1
4	3	1
5	3	2
6	4	2
7	4	3
8	5	3
9	5	4

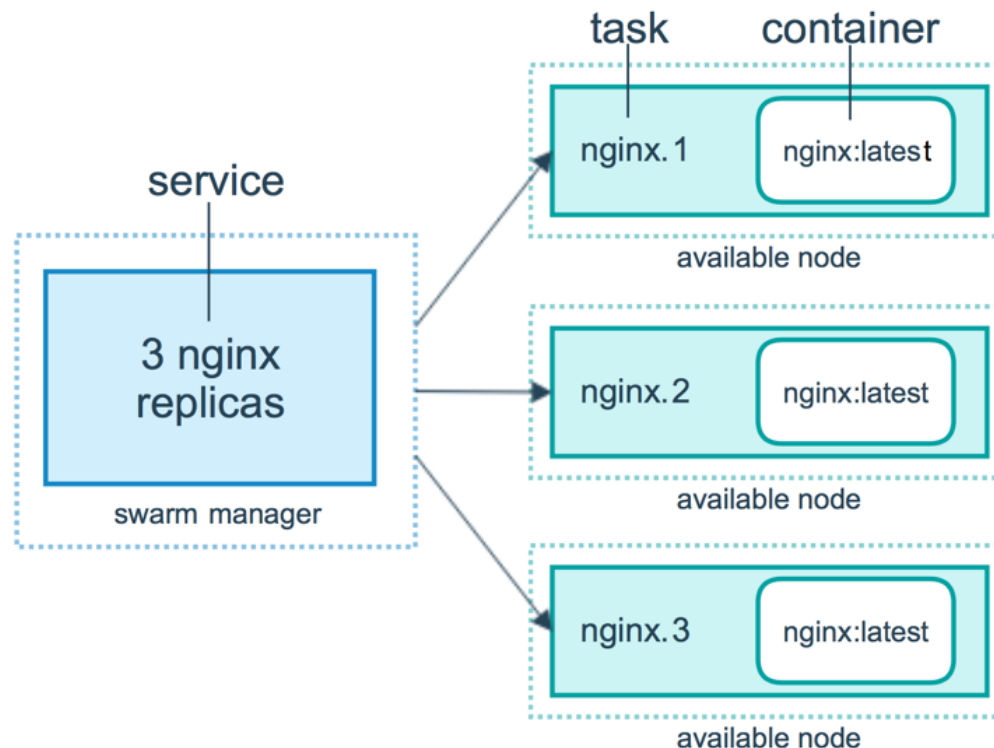
See also: [Docker swarm admin guide](#), [Raft consensus algorithm](#)

Docker swarm: Services

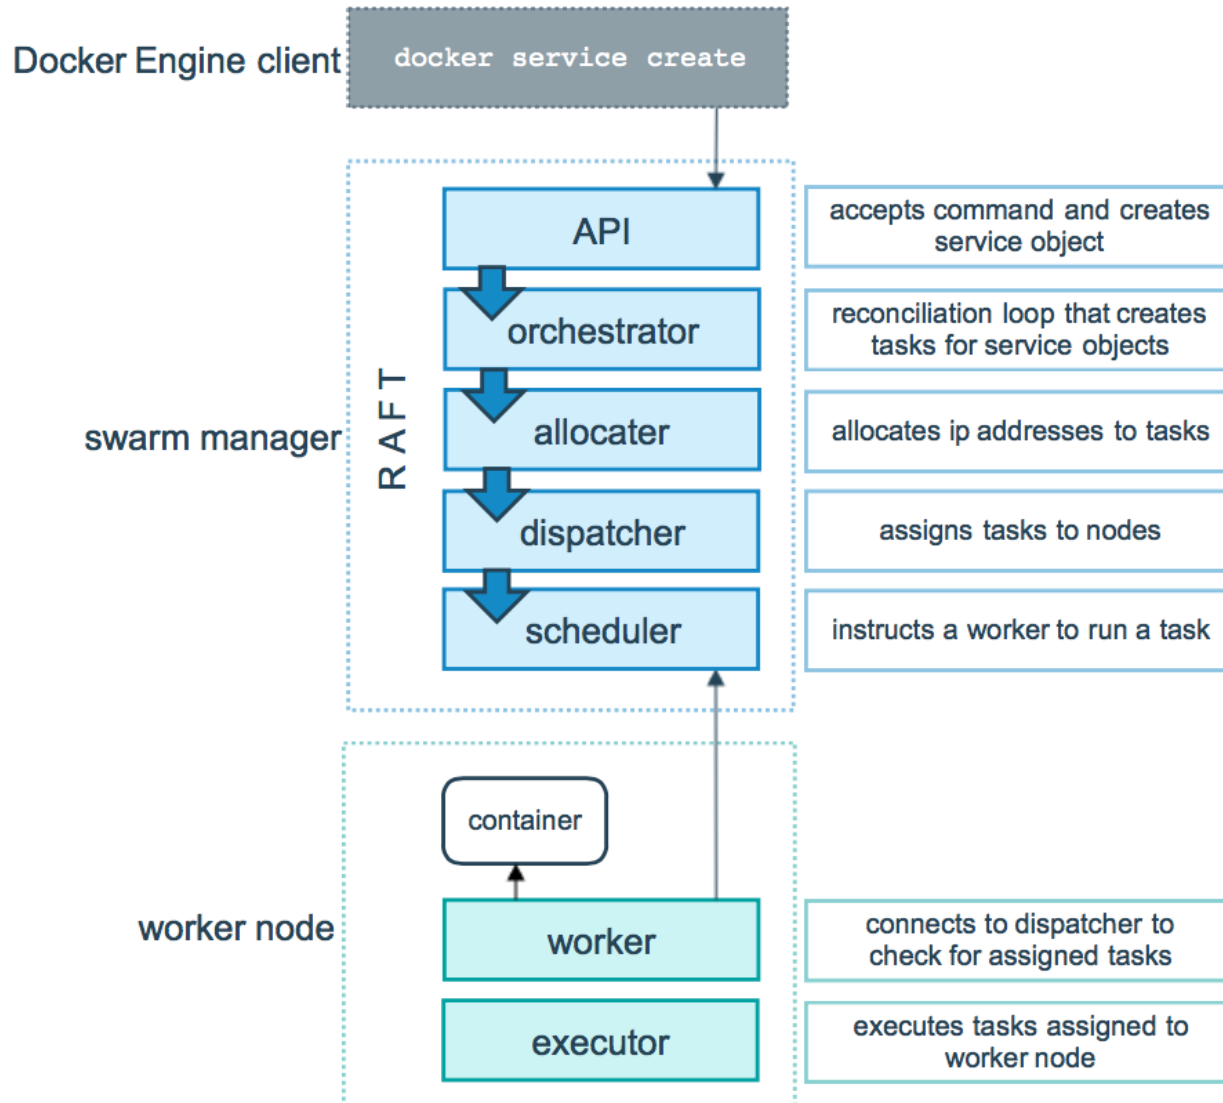
- A service your cluster provides
 - Typically a container that can scale
- Input parameters
 - Mode: Replicated vs. global
 - Name
 - Network
 - Publish (external port)
 - Container id

Docker swarm: Services

- A service your cluster provides
 - Typically a container that can scale



Docker swarm: Services



docker service

- **Run this on your manager node!**
- Equivalent to `docker run`
- Input parameters
 - Mode: Replicated (default) vs. global
 - Name
 - Network
 - Publish (external port)
 - Container id

docker service create

```
--name webservers  
--publish 80:80 nginx
```

docker service

- **Run this on your manager node!**
- `docker service ls`
- `docker service ps webserver`
- `docker service inspect webserver`

Docker swarm: Scaling

- The beauty of docker
 - Containerize once, run anywhere!
- `docker service scale webserver=5`
- `docker service ls`
- `docker service inspect webserver`

See also: [Scaling services](#)

The problem with bottlenecks

- Single point of entry
 - Content delivery network
- Single point of failure
 - Redundancy
- Congestion
 - Scaling!

Recap

- Docker-machine
- Docker swarm
- Docker service

- Coming up:
 - Upgrade strategies
 - Adding/removing nodes
 - Load balancing
 - Service discovery

Upgrade strategies

- Blue-green
 - Two identical environments
 - Only one is live at any time
- Canary
 - Deploy to a small group first, then deploy to the rest
- Rolling
 - Deploy in rolling iterations

See also: [Colorful deployments](#)

Docker swarm: Rolling updates

- What Docker swarm does
 - 1) Stop the first task
 - 2) Schedule update for the stopped task
 - 3) Start the container for the updated task
 - 4) If the update to a task returns RUNNING, wait for the specified delay period then start the next task
 - 5) If, at any time during the update, a task returns FAILED, pause the update
- PS: You need at least two replicas!

See also: [Rolling updates](#)

Docker swarm: Updating

- The beauty of docker
 - Containerize once, run anywhere!
 - In different versions
- `docker service update`
 - `--image nginx:1.12 webserver`
- ... or in this case downgrading

See also: [Scaling services](#)

Docker swarm: Adding nodes

- Two types:
 - Managers
 - Workers
- Security
 - You shouldn't be able to join anywhere
 - Solution: tokens (worker/manager)
- `docker swarm join --token X IP`

See also: [Adding nodes](#)

Docker swarm: Removing nodes

- Two types:

- Managers
- Workers

- From the worker:

```
docker swarm leave
```

- From the manager:

```
docker node update --availability drain <node>  
docker swarm rm
```

See also: [Draining nodes](#)

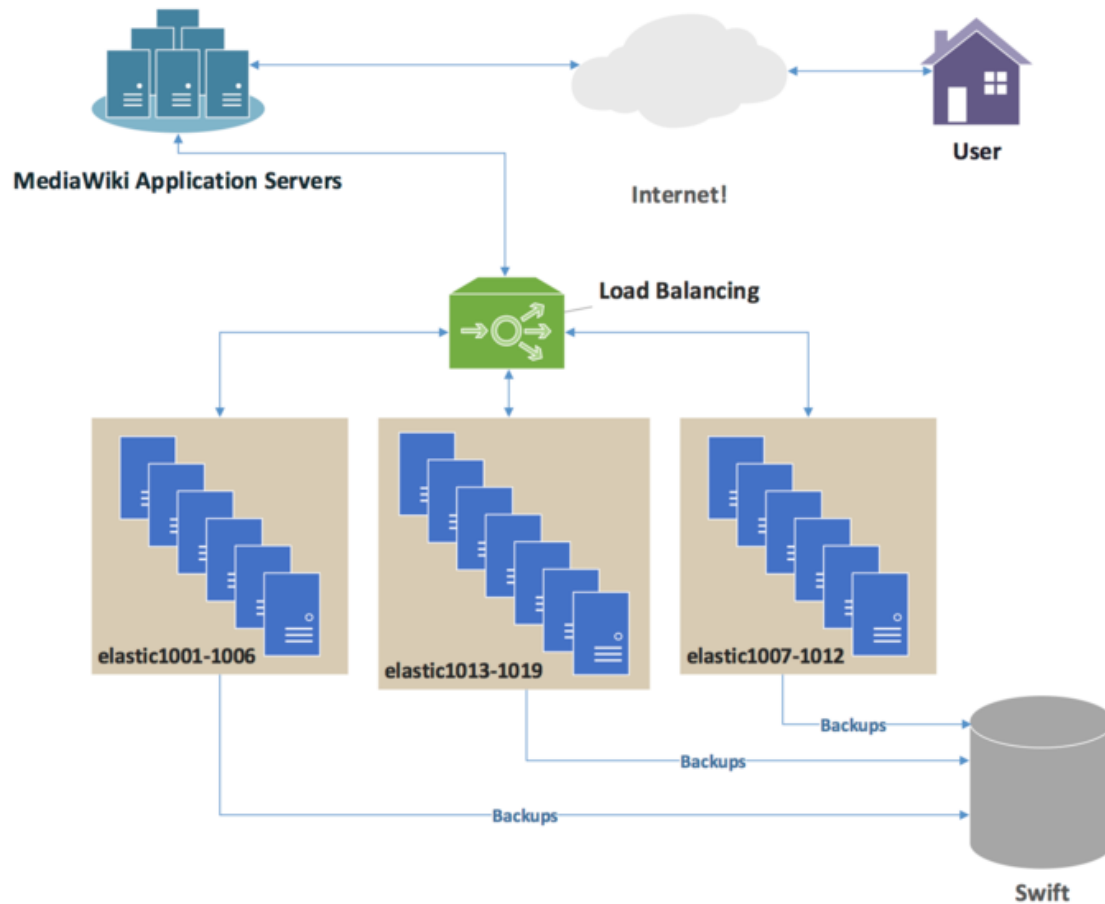
Recap

- Docker-machine
- Docker swarm
 - Creating a swarm
 - Adding/removing nodes
- Docker service
 - Upgrading services
- Coming up:
 - Load balancing
 - Service discovery

The problem with bottlenecks

- Single point of entry
 - Content delivery network
- Single point of failure
 - Redundancy
- Congestion
 - Scaling!

Load balancing



Load balancing strategies

1) Round robin delegation

- Delegate to servers one by one

2) DNS delegation

- Delegate by DNS zones
- Distribute geographically

3) Client-side load balancing

- Clients have a list of servers
- Choose one randomly

See also: [Load balancing on Wikipedia](#)

Load balancing scenarios

- Peak capacity
 - You're Facebook, users gets home from work
- Reduced capacity
 - You're Facebook, users are sleeping
- DDOS

Other load balancing strategies

- Caching
- Compression (less data)
- Multi-threading
- Blocking
 - IP regions
 - Single targets


Load balancing software

- Nginx, haproxy, apache server
 - Many types, same purpose
- “Reverse proxy”
 - Not a proxy for connections out but connections in

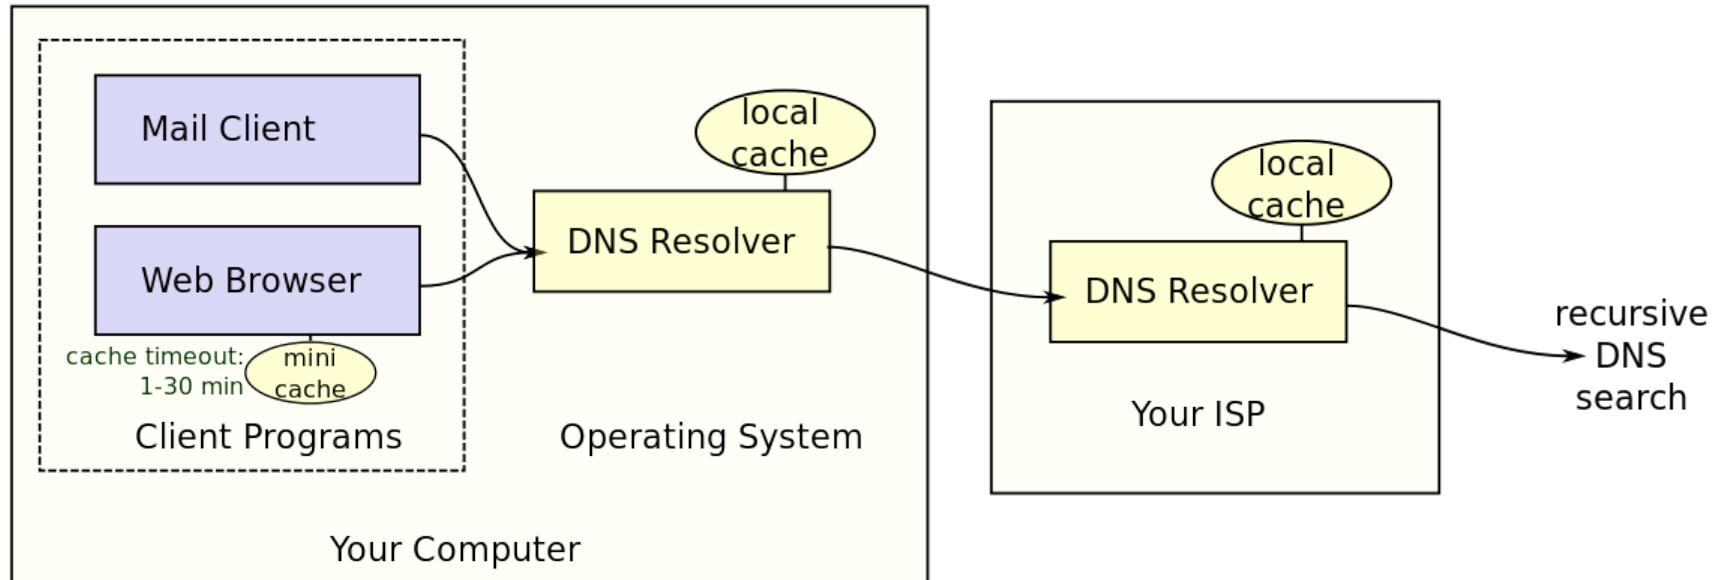
Domain name

- Read from right to left
- Fully qualified domain name (FQDN)
 - Begins with the root domain "."

- 1) Root domain
- 2) Top-level domain
- 3) Second or third level



DNS Resolving



DNS scaling

- Domain name system
 - IP addresses cached on DNS servers
- DNS A record with multiple entries:
 - 1) Request: 0.0.0.1
 - 2) Request: 0.0.0.2
 - 3) Request: 0.0.0.3...
- Your DNS provider can route your traffic
 - Example: DDOS

See also: [DNS on wikipedia](#), [List of DNS record types](#)



DNS scaling

- Domain name system
 - IP addresses cached on DNS servers
- DNS A record with multiple entries:
 - 1) Request: 0.0.0.1
 - 2) Request: 0.0.0.2
 - 3) Request: 0.0.0.3...
- Your DNS provider can route your traffic
 - Example: DDOS

See also: [DNS on wikipedia](#), [List of DNS record types](#)

Content delivery network (CDN)

- “Geographically distributed network of proxy servers and their data centers”



See also: [CDN on Wikipedia](#)

Load balancing strategies

1) Round robin delegation

- Delegate to servers one by one

2) DNS delegation

- Delegate by DNS zones
- Distribute geographically

3) Client-side load balancing

- Clients have a list of servers
- Choose one randomly

See also: [Load balancing on Wikipedia](#)

Docker swarm load balancing

- By default: ingress network
- The ingress network is a special overlay network that facilitates load balancing among a service's nodes.
- When any swarm node receives a request on a published port, it hands that request off to a module called IPVS. IPVS keeps track of all the IP addresses participating in that service, selects one of them, and routes the request to it, over the ingress network.
- The ingress network is created automatically when you initialize or join a swarm.

See also: [Docker swarm networking](#)

Recap

- Docker-machine
- Docker swarm
 - Creating a swarm
 - Adding/removing nodes
- Docker service
 - Upgrading services
- Docker network
 - Ingress: load balancing
- Coming up:
 - Service discovery

Service discovery

- Automatically discover machines providing the same service
- DNS A record with multiple entries:
 - 1) Request: 0.0.0.1
 - 2) Request: 0.0.0.2
 - 3) Request: 0.0.0.3...

See also: [DNS-SD on Wikipedia](#), [RFC2782](#)

Service discovery in docker

- Overlay networks
 - manage communications among the Docker daemons participating in the swarm

```
docker network create
```

```
--driver overlay monitoring
```

```
nslookup tasks.docker-exporter
```

See also: [Docker swarm networking](#)

Monitoring via docker-machine

- Docker-machine experimental feature
 - Inbuilt prometheus monitoring
- `docker-machine create`
 - `--driver virtualbox`
 - `--engine-opt experimental`
 - `--engine-opt metrics-addr=0.0.0.0:4999 mybox`

See also: [Docker metrics in Prometheus](#)

Monitoring via docker-machine

- Docker-machine experimental feature
 - Inbuilt prometheus monitoring
- One small problem... We have to expose it on a specific docker network

```
docker
  service create
    --mode global
    --name docker-exporter
    --network monitoring
    --publish 4999
    -e IN=172.18.0.1:4999
    basi/socat:v0.1.0
```

See also: [Docker metrics in Prometheus](#)

A note on configuration

- No longer a common file system
- How to handle configuration?

```
docker config create prometheus-config prometheus.yml
```

```
docker service create --name prometheus  
    --config=prometheus-config,target=/etc/prometheus/prometheus.yml
```

See also: [Docker metrics in Prometheus](#)

Monitoring via docker-machine

- Putting it all together
- We have docker-exporters and a prometheus configuration file to listen for the dns service discovery

```
docker service create
```

```
--name prometheus -p 9090:9090  
--config src=prometheus,  
        target=/etc/prometheus/prometheus.yml  
--network monitoring prom/prometheus
```

See also: [Docker metrics in Prometheus](#)

Recap

- Docker-machine
- Docker swarm
 - Creating a swarm
 - Adding/removing nodes
- Docker service
 - Upgrading services
- Docker network
 - Ingress: load balancing
- Service discovery
 - Prometheus monitoring via DNS service discovery

Next hand-in

Deadline: 22nd of November 23:59:55

1) Create a docker swarm with at least 3 vms

2) Create your services in docker swarm

1) At least three: Your service, Prometheus and logging (ELK)

3) Optional: Add monitoring to your docker swarm

1) Create an overlay network

2) Create the socat global service

3) Reconfigure Prometheus to use DNS service discovery

- Remember to put Prometheus on the same network!

Next hand-in

Deadline: **22nd of November 23:59:55**

Hand-in:

1) Four lines of text describing

- What Docker swarm is and why we need it
- Why it can help us to eliminate bottlenecks

2) Printout of:

- `docker node ls`
- `docker service ls`